



MODELS AND ALGORITHMS FOR INTEGRATION OF VEHICLE AND CREW SCHEDULING

RICHARD FRELING[†], DENNIS HUISMAN^{*}, AND ALBERT P. M. WAGELMANS

*Erasmus Center for Optimization in Public Transport (ECOPT) and Econometric Institute,
Erasmus University Rotterdam, The Netherlands*

ABSTRACT

This paper deals with models, relaxations, and algorithms for an integrated approach to vehicle and crew scheduling for an urban mass transit system with a single depot. We discuss potential benefits of integration and provide an overview of the literature which considers mainly partial integration. Our approach is new in the sense that we can tackle integrated vehicle and crew scheduling problems of practical size.

We propose new mathematical formulations for integrated vehicle and crew scheduling problems and we discuss corresponding Lagrangian relaxations and Lagrangian heuristics. To solve the Lagrangian relaxations, we use column generation applied to set partitioning type of models. The paper is concluded with a computational study using real life data, which shows the applicability of the proposed techniques to practical problems. Furthermore, we also address the effectiveness of integration in different situations.

KEY WORDS: vehicle scheduling; crew scheduling; integrated planning; column generation; Lagrangian relaxation

1. INTRODUCTION

Although in the early 1980s several researchers recognized the need to integrate vehicle and crew scheduling for an urban mass transit system, most of the algorithms published in the literature still follow the sequential approach where vehicles are scheduled before, and independently of, crews. Algorithms incorporated in commercially successful computer packages use this sequential approach as well, while sometimes integration is dealt with at the user level (e.g., see Darby-Dowman et al. (1988)). In the operations research literature, only a few publications address a simultaneous approach to vehicle and crew scheduling. None of those publications makes a comparison between simultaneous and sequential scheduling. Hence, they do not provide any indication of the benefit of a simultaneous approach. Only in the airline literature, there recently appeared a publication about the benefits of an integrated approach between aircraft routing (which is equivalent to vehicle scheduling) and crew scheduling. Klabjan et al. (2002) showed that partial integration of schedule planning, aircraft routing, and airline crew scheduling can improve the quality of the crew scheduling solution enormously. To be more

^{*} Correspondence to: Dennis Huisman, Econometric Institute, Erasmus University Rotterdam, P.O. Box 1738, 3000 DR Rotterdam, The Netherlands. E-mail: huisman@few.eur.nl

[†] In memoriam: Richard Freling passed away on January 29, 2000 at the age of 34.

specific, the flight time credit (FTC), which is defined as the percentage of excess cost above flying, is improved by a factor of two.

In this paper, we consider a complete integration of vehicle and crew scheduling for the single depot case. An overview of the potential benefits of integration is provided in Freling Wagelmans, and Paixão (1999). To evaluate the effectiveness of this approach, we also consider the traditional sequential approach and the opposite sequential approach of scheduling crews before, and independently of, vehicles.

The paper is organized as follows. Section 2 deals with problem definitions. An overview of the literature on simultaneous vehicle and crew scheduling approaches is provided in Section 3.

All algorithms proposed in this paper rely on column generation applied to set partitioning type of models (e.g., see Freling (1997) and Desrochers and Soumis (1989)). We use Lagrangian relaxations and Lagrangian heuristics with column generation (Section 4) and we discuss quality guarantees for both the integration and independent crew scheduling. The column generation (pricing) problem allows us to exploit the network flow structure of the Crew Scheduling Problem (CSP), and is decomposed in two or three phases, each consisting of (constrained) shortest path type of problems.

In Sections 5–8, we propose mathematical formulations and algorithms for crew scheduling, complete integration of vehicle and crew scheduling, where a crew is allowed to work on more than one vehicle during a duty, for the situation where this is not allowed, and for independent crew scheduling, respectively. The paper is concluded with a computational study using real life data from a mass transit company in the Netherlands (Section 9).

2. PROBLEM DEFINITION

The *vehicle and crew scheduling problem* (VCSP) is the following: Given a set of service requirements or *trips* within a fixed planning horizon, find a minimum cost schedule for the vehicles and the crews, such that both the vehicle and the crew schedules are feasible and mutually compatible. Each trip has fixed starting and ending times, and the traveling times between all pairs of locations are known. A vehicle schedule is feasible if (1) each trip is assigned to a vehicle, and (2) each vehicle performs a feasible sequence of trips, where a sequence of trips is feasible if it is feasible for a vehicle to execute each pair of consecutive trips in the sequence. From a vehicle schedule it follows which trips have to be performed by the same vehicle and this defines so-called vehicle *blocks*. The blocks are subdivided at *relief points*, defined by location and time, where and when a change of driver may occur. A *task* is defined by two consecutive relief points and represents the minimum portion of work that can be assigned to a crew. These tasks have to be assigned to crew members. The tasks that are assigned to the same crew member define a crew *duty*. Together the duties constitute a crew schedule. Such a schedule is feasible if (1) each task is assigned to one duty, and (2) each duty is a sequence of tasks that can be performed by a single crew, both from a physical and a legal point of view. In particular, each duty must satisfy several complicating constraints corresponding to work load regulations for crews. Typical examples of such constraints are maximum working time without a break, minimum break duration, maximum total working time, and maximum duration. The cost of a duty is usually a combination of fixed costs such as wages, and variable costs such as overtime payment. Finally, we define a *piece (of work)* as a sequence of tasks on one vehicle block without a break that can be performed by a single crew member without interruption.

We make the following assumptions:

1. There is only one depot, all vehicles are available at all time, and they are identical. With these characteristics, the *single depot vehicle scheduling problem* (SDVSP), in which one has to construct feasible vehicle schedules that minimize the total vehicle costs, is polynomially solvable (e.g., see Freling, Paixão, and Wagelmans (2001)).
2. The cost function for the VCSP is the summation of the vehicle and crew scheduling cost functions, where the primary vehicle scheduling objective is to minimize the number of vehicles, while the primary crew scheduling objective is to minimize the number of crews. An additional term can be added to the cost function which corresponds to variable vehicle costs (e.g., distance travelled) and variable crew costs (e.g., overtime payment).
3. The feasibility of a piece only depends on its duration, which is limited by a minimum and maximum piece length. This follows from legal regulations.

The three assumptions make our approach in principle applicable to bus and driver scheduling.

We distinguish between two types of tasks, namely *trip tasks* corresponding to (parts of) trips and *dh-tasks* corresponding to deadheading. A *deadhead* is a period that a vehicle is moving to or from the depot, or a period between two trips that a vehicle is outside of the depot (possibly moving without passengers). All trip tasks need to be covered by a crew, while the covering of dh-tasks depends on the vehicle schedules and determines the compatibility between vehicle and crew schedules. In particular, each dh-task needs to be assigned to a crew if and only if its corresponding deadhead is assigned to a vehicle. Note that more than one trip task may correspond to a single trip, depending on the relief points along that trip. Similarly, more than one dh-task may correspond to a single deadhead. For example, a deadhead between the end location of trip i and the starting location of trip j which passes by the depot corresponds to two dh-tasks, one from the end location of trip i to the depot and the other from the depot to the starting location of trip j . Here we assume that waiting at the depot is not a task because vehicle attendance at the depot is not necessary.

3. LITERATURE REVIEW

The traditional sequential strategy is strongly criticized by Bodin et al. (1983). This is motivated by the fact that in North American mass transit organizations the crew costs dominate vehicle operating costs, and in some cases reach as high as 80% of total operating costs. Although simultaneous vehicle and crew scheduling is of significant practical interest, only a few approaches of this kind have been proposed in the literature. They mainly deal with bus and driver scheduling and fall into one of the following three categories:

1. Scheduling of vehicles during a heuristic approach to crew scheduling.
2. Inclusion of crew considerations in the vehicle scheduling process; the actual crew scheduling is only carried out afterwards.
3. Complete integration of vehicle and crew scheduling.

Most of the approaches are of the first category and are based on a heuristic procedure proposed by Ball, Bodin, and Dial (1983). This procedure involves the definition of a scheduling network, which consists of vertices characterized by parts of trips called *d-trips* that have to be executed by one vehicle and crew, and two vertices s and t representing the depot. The arcs can

be grouped into two categories, those which indicate that a crew and vehicle proceed from one d-trip to another and those which indicate that only the crew proceeds from one d-trip to another (*crew-only* arcs). The solution procedure is decomposed into three components, emphasizing the CSP: A piece construction component, a piece improvement component, and a duty generation component. The piece construction routine generates a set of pieces whose time duration is less than some constant T . (Note that this corresponds to vehicle scheduling with time constraint.) This problem is solved by a heuristic based on matching. In the second step pairs of short pieces are combined into partial duties, while in a third step pairs of these duties and longer pieces are combined into two- and three-piece duties. The second step uses a matching based interchange heuristic and the third one can be solved as a matching problem. Vehicle schedules are generated simultaneously by deleting the crew-only arcs and fixing arcs used by pieces in the solution. This procedure is applied to large size VCSP instances which correspond to the entire physical network, that is, all lines are considered at once, while no restrictions are placed on interlining, that is, a crew may work on an arbitrary number of lines.

Other heuristic approaches of the first category are proposed by Tosini and Vercellis (1988), Falkner and Ryan (1992), and Patrikalakis and Xerocostas (1992). All these approaches use a similar crew scheduling network as in Ball, Bodin, and Dial (1983).

Approaches of the second category are proposed by Darby-Dowman et al. (1988) as an interactive part of a decision support system, and by Scott (1985) who heuristically determines vehicle schedules which take crew costs into account. An initial vehicle schedule is heuristically modified according to estimated marginal costs associated with a small change in the current vehicle schedule. The estimated marginal costs are obtained by solving the linear programming dual of the HASTUS crew scheduling model (see Rousseau and Blais (1985)). Results obtained with public transport scheduling problems in Montréal, show a slight decrease in estimated crew costs.

Only very recently, approaches of the third category, that is, a complete integration of vehicle and crew scheduling, have been proposed. The first mathematical formulation was proposed by Freling, Boender, and Paixão (1995). This formulation is similar to the one considered in this paper (see also Freling, Wagelmans, and Paixão (1999)). Haase and Friberg (1999) propose an exact algorithm for the VCSP. Both the vehicle and crew scheduling aspects are modelled by using set partitioning type of constraints. A *branch-and-cut-and-price* algorithm is proposed, that is, column generation and cut generation are combined in a branch-and-bound algorithm. The column generation master problem corresponds to the linear programming (LP) relaxation, while the pricing problem corresponds to a shortest path problem for generating vehicle schedules and a constrained shortest path problem for generating duties. The constrained shortest path problem is solved using the algorithm proposed by Desrochers (1988). Preliminary computational results indicate that problems with 10 trips can be solved within 1 min, problems with 20 trips can be solved in more than an hour, and problems with 30 trips cannot be solved.

Haase, Desaulniers, and Desrosiers (2001) propose an approach which solves a CSP that incorporates side constraints for the buses. This is done in such a way that the solution of this problem guarantees that an overall optimal solution is found after constructing a compatible vehicle schedule. The solution approach is based on a multicommodity network flow formulation for the CSP with side constraints, which is solved by a branch-and-price algorithm. For larger problem instances a heuristic version of the algorithm is used. Computational experiments with random data instances, simulating an urban mass transit environment, show that instances with up to 300 trips can be solved within 94 min of CPU time (on a SUN ULTRA-10/440 workstation) and with an integrality gap of less than 1.1%. Furthermore, out of

10 instances with 150 trips, 6 instances could be solved to optimality within 3 h of CPU time. We note, however, that their computational experiments cannot be compared in a straightforward way with ours, since there are some important differences between their random data and our real life data (see Section 9).

Gaffi and Nonato (1999) propose a heuristic algorithm based on Lagrangian relaxation with column generation for an integrated approach to multiple-depot vehicle scheduling and crew scheduling. They also propose a mathematical formulation similar to the formulation presented in this paper, with some modifications for the multiple-depot case. Their approach is developed for the ex-urban mass transit setting, where crews are tightly dependent on the vehicle activities or dead-heading of crew is highly constrained. Computational results are provided for Italian public transit operators, which show some improvements over the results of a sequential approach. CPU times (on a Power PC 604, 180 Mhz) are over 24 h for cases with up to 257 trips, and on average 2–6 h.

Finally, Klabjan et al. (2002) as mentioned in the introduction, recently investigated a different type of approach for integrated fleet planning and crew scheduling. They incorporate plane count constraints in their CSP, which is essentially the same as the side constraints in Haase, Desaulniers, and Desrosiers (2001). They use a LP based branch-and-bound methodology.

4. COLUMN GENERATION IN COMBINATION WITH LAGRANGIAN RELAXATION

We propose Lagrangian heuristics with column generation for obtaining feasible solutions with quality guarantee for the different models. We assume that the reader is already familiar with the techniques of Lagrangian relaxation and subgradient optimization (surveys on these topics are Geoffrion (1974), Fisher (1981), and Beasley (1995)).

Column generation is a technique that is used for problems with a huge number of variables (e.g., see Barnhart et al. (1998)). The general idea is to solve a sequence of reduced problems, where each reduced problem contains only a small portion of the set of variables (columns). After solving a reduced problem a new set of columns is obtained by using dual information of the solution. The column generation algorithm converges once it is established that the optimal solution based in the current set of columns cannot be improved upon by adding more columns. Then the optimal solution of the reduced problem is the optimal solution of the overall problem. We will refer to the reduced problem as the *master problem* and to the problem of generating a new set of columns as the *pricing problem*.

Usually column generation is used in the context of linear programming, but here we will use it in combination with Lagrangian relaxation (see also Freling (1997) and Carraresi, Girardi, and Nonato (1995)) to compute lower bounds on the optimal solution. Furthermore, we will use the columns which are generated to compute the lower bound to construct a feasible solution. This approach has been chosen to solve the different models for the following reasons:

1. For these problems, a set of columns generated to compute the lower bound turns out to be a set from which we can select a reasonably good feasible solution;
2. Since we compute a lower bound on the optimal solution, we obtain an indication about the quality of the constructed feasible solution;
3. Lagrangian relaxation has shown to provide tight bounds for set partitioning type of problems (see e.g., Beasley (1995));

4. Because of the number of constraints, using a linear programming relaxation is not a realistic option for the VCSP1 model, which will be discussed in Section 6.

We also mention that Carraraesi, Girardi, and Nonato (1995) proposed a similar algorithm for the Crew Scheduling Problem (CSP). The only difference between their and our approach is that in subsequent reduced problems they replace the set of columns by a completely new one, while we only add columns.

5. CREW SCHEDULING

In this section we discuss the CSP, which is part of the sequential approach. However, before we can solve the CSP, we have to solve the vehicle scheduling problem (VSP) and generate all feasible pieces and (optional) all feasible duties first (see Section 5.3). To solve the VSP, we use a very efficient auction algorithm, which is described in detail in Freling, Paixão, and Wagelmans (2001). The solution of the VSP gives a set of vehicle blocks on which we can define the relief points. This results in a set of tasks from which we can easily enumerate all feasible pieces, since a piece is a feasible sequence of consecutive tasks on the same vehicle block only restricted by its duration.

5.1. Mathematical formulation

In this subsection we give a mathematical formulation for the CSP. Let d_k be the cost of duty $k \in K$, where K is the set of all feasible duties, and we define $K(i) \subseteq K$ as the set of duties covering task $i \in I$. Consider binary decision variable x_k indicating whether duty k is selected in the solution or not. In the *set covering formulation* of the CSP, the objective is to select a minimum cost set of feasible duties such that each task is included in at least one of these duties. This is the following 0–1 linear program:

$$\min \sum_{k \in K} d_k x_k \tag{1}$$

$$\sum_{k \in K(i)} x_k \geq 1 \quad \forall i \in I \tag{2}$$

$$x_k \in \{0, 1\} \quad \forall k \in K \tag{3}$$

The advantage of working with this formulation instead of a *set partitioning* one is that it is easier to solve (a similar remark can be made for relaxations of the respective problems). After solving the set covering formulation, we can always change the solution into a set partitioning solution by deleting over-covers of trips. In fact, this boils down to changing some of the selected duties. Instead of being the driver, the person who is assigned to such a duty will make the trip as a passenger. Note that such changes affect neither the feasibility nor the cost of the duties involved.

5.2. Algorithm

As mentioned before, we solve the CSP as a set covering problem. A global description of the Lagrangian heuristic is given in Figure 1.

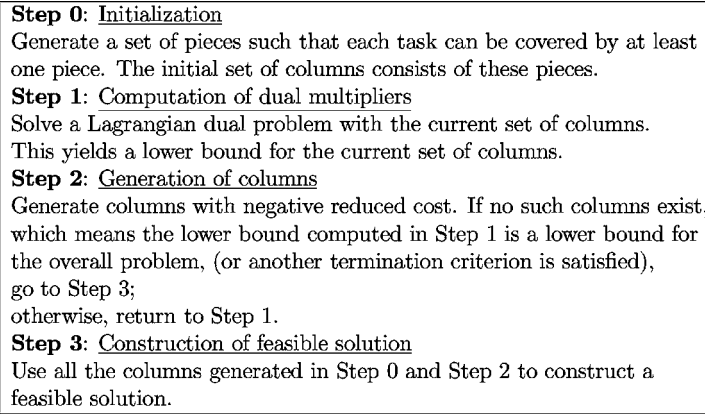


Figure 1. Solution method for CSP

Suppose that, at some point, K is the set of columns (duties) that we are considering in the master problem. We compute the lower bound with respect to these columns as follows. We associate nonnegative Lagrangian multipliers $\lambda_1, \lambda_2, \dots, \lambda_{|I|}$ with the constraints (2) and then the remaining Lagrangian subproblem amounts to *pricing out* the x variables. We use subgradient optimization to solve the Lagrangian dual problem approximately. Besides a good lower bound, we also obtain a final set of Lagrange multipliers. These are used to compute the reduced cost \bar{d}_k of columns $k \notin K$ in the column generation step, where $\bar{d}_k = d_k - \sum_{i \in I(k)} \lambda_i$ with respect to the Lagrangian multipliers $\lambda = (\lambda_1, \lambda_2, \dots, \lambda_{|I|})$, and $I(k)$ is the set of tasks in duty k . We add a number of duties with negative reduced cost to K and repeat the subgradient optimization. Actually, to assure that we do not generate columns twice, we modify the Lagrange multipliers before we generate new columns. The modifications are such that all columns already in K get nonnegative reduced cost, while not decreasing the value of the Lagrangian function. For details on the greedy heuristic that we use, we refer to Freling (1997) and Carraraesi, Girardi, and Nonato (1995).

We stop if there are no duties left with negative reduced cost or if the lower bound obtained by the subgradient optimization does not significantly change for a number of iterations. In the first case we obtain a true lower bound (although not necessarily the best Lagrangian lower bound); in the second case we do not have this guarantee.

Finally we compute a feasible solution by solving a set covering problem in which we consider all the columns which have been generated along the way. We can either do this exactly, using a general or specialized integer programming solver, or heuristically. Most often we use the set covering heuristic of Caprara, Fischetti, and Toth (1999) with a few changes proposed in Freling (1997).

5.3. Generation of duties

Duties consist of a number of pieces with a given maximum number of pieces. In practice this maximum is very often equal to 2 or 3. This is the reason why we simply enumerate all possible combinations of pieces and check if such a combination is feasible, until we find a specified number of duties with, negative reduced costs. (To generate duties consisting of more than

3 pieces, one may use a duty generation network, as proposed by Freling (1997).) The reduced cost of a duty can be easily computed if the reduced cost of a piece is already known: the reduced cost of a duty is equal to the sum of the reduced cost of the pieces it is built from, plus the reduced cost of the breaks between the pieces.

6. COMPLETE INTEGRATION: GENERAL CASE

In this section, we propose a mathematical formulation for the VCSP in the general case. We deal separately with the situation where changeovers are not allowed (see Section 7). A *changeover* is the change of a vehicle of a bus driver during his break.

6.1. Mathematical formulation

Patrikalakis and Xerocostas (1992) presented the first mathematical formulation, for the VCSP. This is a simplified version of the model presented here. Our formulation is valid under the assumptions that short arcs correspond to only one dh-task and *continuous attendance* is required, that is, there is always a driver present if the bus is outside the depot. If these assumptions are not met, we can use a slightly different formulation; for details we refer to Huisman (2000).

The mathematical formulation we propose for the VCSP is a combination of the quasi-assignment formulation for the vehicle scheduling problem, and the set partitioning formulation for crew scheduling discussed in Section 5.1. The quasi-assignment part assures the feasibility of vehicle schedules, while the set partitioning part assures that each trip task is assigned to a duty and each dh-task is assigned to a duty if and only if its corresponding deadhead is part of the vehicle schedule. Before providing the mathematical formulation, we need to recall and to introduce some notation.

Let $N = \{1, 2, \dots, n\}$ be the set of trips, numbered according to increasing starting time, and let $E = \{(i, j) \mid i < j, i, j \text{ compatible}, i \in N, j \in N\}$ be the set of deadheads. Two trips are compatible if they can be driven by the same vehicle after each other. Let s and t both represent the depot at location d . We define the vehicle scheduling network $G = (V, A)$, which is an acyclic directed network with nodes $V = N \cup \{s, t\}$, and arcs $A = E \cup (s \times N) \cup (N \times t)$. A path from s to t in the network represents a feasible schedule for one vehicle, and a complete feasible vehicle schedule is a set of disjoint paths from s to t such that each node in N is covered. Let c_{ij} be the vehicle cost of arc $(i, j) \in A$, which is usually some function of travel and idle time. Furthermore, a fixed cost for using a vehicle can be added to the cost of arcs (s, i) or (j, t) for all $i, j \in N$.

To reduce the number of constraints, we assume that a vehicle returns to the depot if it has an idle time between two consecutive trips which is long enough to let it return. In that case the arc between the trips is called a *long arc*; the other arcs between trips are called *short arcs*. Denote $A^s \subset A$ and $A^l \subset A$ as the sets of short and long arcs, respectively.

As before, K denotes the set of duties and d_k denotes the crew cost of duty $k \in K$, respectively. We assume that deadheads to and from the depot correspond to one dh-task each. Suppose $(i, j) \in A^l$ and let el_i and bl_j denote the ending and starting location of trips i and j , respectively. Then we let $K(i, t)$ and $K(s, j)$ denote the set of duties covering the dh-task from el_i to the depot and from the depot to bl_j , respectively. Furthermore, I_1 denotes the set of trip tasks and $K(p)$ is the set of duties covering trip task $p \in I_1$, where we assume that a trip corresponds to one task. Note that this means that there is a one-to-one correspondence between trips and trip tasks.

$K(i, j)$ denotes the set of duties covering dh-tasks corresponding to deadhead $(i, j) \in A^s$. Decision variables x_k are defined as before, that is, x_k indicates whether duty k is selected in the solution or not. Furthermore, there are decision variables y_{ij} that indicate whether a vehicle covers trip j immediately after trip i or not. The VCSP can be formulated as follows.

VCSP1:

$$\min \sum_{(i,j) \in A} c_{ij} y_{ij} + \sum_{k \in K} d_k x_k \quad (4)$$

$$\sum_{\{j: (i,j) \in A\}} y_{ij} = 1 \quad \forall i \in N \quad (5)$$

$$\sum_{\{j: (i,j) \in A\}} y_{ij} = 1 \quad \forall j \in N \quad (6)$$

$$\sum_{k \in K(p)} x_k = 1 \quad \forall p \in I_1 \quad (7)$$

$$\sum_{k \in K(i,j)} x_k - y_{ij} = 0 \quad \forall (i,j) \in A^s \quad (8)$$

$$\sum_{k \in K(i,t)} x_k - y_{it} - \sum_{\{j: (i,j) \in A'\}} y_{ij} = 0 \quad \forall i \in N \quad (9)$$

$$\sum_{k \in K(s,j)} x_k - y_{sj} - \sum_{\{j: (i,j) \in A'\}} y_{ij} = 0 \quad \forall j \in N \quad (10)$$

$$x_k, y_{ij} \in \{0, 1\} \quad \forall k \in K, \forall (i,j) \in A \quad (11)$$

The objective is to minimize the sum of total vehicle and crew costs. The first two sets of constraints, (5) and (6), correspond to the quasi-assignment formulation for the SDVSP. They assure that each trip is assigned to exactly one predecessor and successor, that is, these constraints guarantee that the network is partitioned into a set of disjoint paths from s to t . Constraints (7) assure that each trip task p will be covered by one duty in the set $K(p)$. Furthermore, constraints (8), (9), and (10) guarantee the link between dh-tasks and deadheads in the solution, where deadheads corresponding to short and long arcs in A are considered separately. In particular, constraints (8) guarantee that each deadhead from i to j is covered by a duty in the set $K(i,j)$ if and only if the corresponding short arc is in the vehicle solution. However, notice that the constraints (8) are correct under the assumptions we made, but under different assumptions we may have to change the equality sign to an “ \geq ” sign (see Huisman (2000)). The other two constraint sets, (9) and (10), ensure that the dh-tasks from el_i to t and from s to bl_j , possibly corresponding to long arc $(i, j) \in A$, are both covered by one duty if and only if the corresponding deadheads are in the solution. Note that the structure of these last three sets of constraints is such that each constraint corresponds to the possible selection of one duty from a large set of duties, where the fact whether or not a duty has to be selected depends on the values of the corresponding y variables. We will refer to these constraints as *partitioning type of constraints*. The model contains $|A| + |K|$ variables and $4|N| + |A^s| + |I_1|$ constraints, which is already a huge number of variables and a very large number of constraints for instances

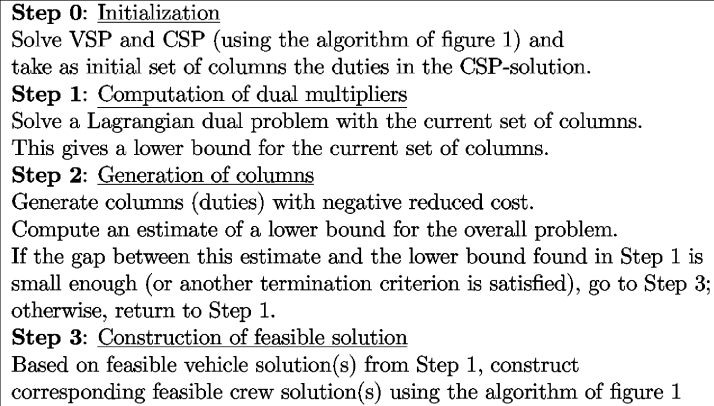


Figure 2. Solution method for VCSP1

with a small number of trips. This is probably the main reason why complete integration has previously received very little attention.

6.2. Algorithm

The algorithm for VCSP1 is shown in Figure 2.

We use the relaxation of model VCSP1, where all set partitioning type of constraints (7)–(10) are first replaced by set covering constraints, which are subsequently relaxed in a Lagrangian way. That is, we associate nonnegative Lagrangian multipliers λ_p , μ_{ij} , ν_i , and ξ_j with constraints (7), (8), (9), and (10), respectively. Then the remaining Lagrangian subproblem can be solved by pricing out the x variables and solving a SDVSP for the y variables. We can use the auction algorithm proposed in Freling, Paixão, and Wagelmans (2001) for solving the SDVSP. Notice that we get a feasible vehicle solution every time we solve the SDVSP.

As before, we need an additional procedure to update the Lagrangian multipliers after solving the Lagrangian relaxation. This is necessary to assure that all duties in the current master problem have nonnegative reduced cost so that these duties will not be generated again in the pricing problem.

In contrast to our approach to the CSP, we cannot generate the complete set of duties before we run the actual optimization algorithm, because we do not have a vehicle solution in advance. Therefore, we have to generate the duties during the process. We do this by generating pieces of work, which we will describe in Section 6.4, and then combining these pieces to generate duties in the same way as for the CSP (see Section 5.3).

In every iteration i we compute an estimate LBT_i of the lower bound for the overall problem. Let LBS_i denote the value of the Lagrangian lower bound in iteration i , then the estimate is computed as

$$\text{LBT}_i = \text{LBS}_i + \sum_{k \in K_i} \bar{d}_k \quad (12)$$

where K_i is the set of duties added in iteration i and \bar{d}_k is the reduced cost of duty $k \in K$, which is defined below.

$$\bar{d}_k = d_k - \sum_{p \in I_1(k)} \lambda_p - \sum_{(i,j) \in A^s(k)} \mu_{ij} - \sum_{i \in N^t(k)} \nu_i - \sum_{j \in N^s(k)} \xi_j \quad (13)$$

where $I_1(k)$, $A^s(k)$, $N^t(k)$, and $N^s(k)$ denote the set of trip tasks I_1 in duty k , the set of short arcs A^s in duty k , the set of trips that have a corresponding task in duty k with the end of this trip as starting location and the depot as ending location and, the set of trips that have a corresponding task in duty k with the depot as starting location and the start of this trip as ending location, respectively.

LBT_i is only a true lower bound for the overall problem if all duties with negative reduced cost have been added to the master problem. Of course, we can stop if LBT_i is equal to LBS_i , but in practice we stop earlier, namely if the relative difference is small and LBT_i is a true lower bound, or if there has not been any significant improvement in LBS_i during a number of iterations.

At the end we compute a feasible crew schedule given the (feasible) vehicle schedule which resulted from solving the last Lagrangian subproblem. We do this by solving the CSP (using the algorithm described in Figure 1). Of course, it is also possible to compute more feasible solutions by solving the CSP not only for the vehicle solution from the last iteration, but also for vehicle solutions which were encountered earlier on. A reason to actually do this could be that the gap between the lower and upper bound is quite large, which is an indication that the upper bound could be improved upon.

6.3. The column generation subproblem

For the VCSP, vehicle blocks are not known and a huge number of feasible pieces of work may exist. Although Ball, Bodin, and Dial (1983) have proposed a duty generation network that does not consider pieces of work explicitly, this approach is not interesting for column generation purposes due to the size of the network. Therefore, we propose a two phase procedure for the column generation pricing problem: In the first phase, a *piece generation network* is used to generate a set of pieces of work which serve as input for the second phase where duties are generated. The second phase is equivalent to generating duties for the CSP.

A disadvantage of the two-phase procedure is that it is possible to generate a duty twice in the second phase, which results in an unnecessary large amount of variables. For example, a duty containing the same tasks can be constructed from two different pieces of work. It is very time consuming to check for each new duty generated if it has been generated before. However, we have implemented an efficient procedure to check for double duties after a fixed number of iterations of our duty generation algorithm. The current list of generated duties is then updated so that it contains no double duties, and the algorithm for generating duties is continued.

6.4. Generation of pieces of work

Recall that we have defined a piece of work as a continuous sequence of trip tasks and dh-tasks corresponding to (a part of) one vehicle block, and that this sequence of tasks is only restricted by its duration.

6.4.1. Network structure

The network for piece generation is an extension of the network G for vehicle scheduling (see Section 6.1). Let a *start point* (*end point*) be defined as the relief point corresponding to the start (end) of a vehicle trip. We define the network $G' = (N', A')$, where nodes correspond to the relief points on each trip, and the source s and the sink t represent the depot. Arcs in A' correspond to dh-tasks and trip tasks. Because every relief point has a unique time point and as we assume that the nodes are numbered according to increasing time, G' is acyclic.

Let b_{ij} be the cost associated with each arc $(i, j) \in A'$. These costs should be defined in such a way that the cost of each path is equal to the cost of the corresponding piece, which in turn is itself derived from the costs of the duties. Notice that this is not possible with each arbitrary cost function. Therefore, we implicitly assume that we have a cost function where this is possible, for example, the case with only fixed crew costs. Then $b_{ij} = 0$ for each arc $(i, j) \in A'$. Recall from Section 6.2 that we associate Lagrangian multipliers $\lambda_p, \mu_{ij}, \nu_i$, and ξ_j with constraints (7), (8), (9), and (10), respectively.

Then, for model VCSP1, the reduced cost is defined as

$$\bar{b}_{ij} = \begin{cases} b_{ij} - \lambda_p, & \text{for each arc } (i, j) \text{ corresponding to a trip task } p \\ b_{ij} - \mu_{ij}, & \text{for each arc } (i, j) \text{ corresponding to a short arc } (i, j) \\ b_{ij} - \xi_r, & \text{for each arc } (i, j) \text{ with } i = s \text{ and } j \text{ the start point of trip } r \\ b_{ij} - \nu_r, & \text{for each arc } (i, j) \text{ with } j = t \text{ and } i \text{ the end point of trip } r \end{cases}$$

Thus, the reduced costs on the arcs are defined such that the reduced cost of a path is equal to the reduced cost of the corresponding piece of work. Let tp_i be the point in time associated with node $i \in N' \setminus \{s, t\}$. Each path $P(u, v)$ between two nodes u and v in network G' corresponds to a feasible piece of work if its duration satisfies the time constraint

$$\text{dur}_{\min} \leq tp_v - tp_u \leq \text{dur}_{\max}$$

where dur_{\min} and dur_{\max} denote the minimum and maximum allowed duration of a piece of work, respectively. The duration of a piece of work starting at s and/or ending at t is determined by incorporating travel times. For example, a path s, u, \dots, v, t corresponds to a feasible piece of work if the inequalities

$$\text{dur}_{\min} \leq tp_v - tp_u + \text{trav}(d, bl_u) + \text{trav}(el_v, d) \leq \text{dur}_{\max}$$

hold, where $\text{trav}(d, bl_u)$ and $\text{trav}(el_v, d)$ denote the travel times (possibly including sign-on and sign-off times) from d to bl_u and from el_v to d , respectively.

6.4.2. An all-pairs shortest path algorithm

In this section, we propose a polynomial all-pairs shortest path algorithm for generating a set of pieces. Such a set is generated by solving a shortest path problem between each pair of nodes in network G' that satisfy the constraints on the piece duration. For all feasible paths $P(u, v) = u, \dots, v$ between nodes $u \neq s$ and $v \neq t$ three additional paths are considered, namely path s, u, \dots, v , path u, \dots, v, t , and path s, u, \dots, v, t . In some applications, all crew breaks need to be assigned to the depot and all crew start and end their duty at the depot. In this case, only paths of the form s, u, \dots, v, t are considered. The algorithm is shown in Figure 3.

The procedure $\text{SHORTESTPATHS}(i, k)$ returns the shortest paths $P^*(i, j)$ from i to $j = i + 1, \dots, k$ in $O(m)$ time because G' is acyclic, where m is the number of arcs in the network G' . The procedure $\text{RETRIEVEPIECES}(i, k)$ is shown in Figure 4.

```

for  $i \in N' \setminus \{s, t\}$  do
begin
    Determine  $k \in \arg \max_{j \in N' \setminus \{s, t\}} \{tp_j - tp_i \leq dur_{max}\}$ .
    SHORTESTPATHS( $i, k$ ).
    RETRIEVEPIECES( $i, k$ ).
end.
    
```

Figure 3. Piece generation algorithm

```

for  $j = i + 1, \dots, k$  do
begin
    if  $tp_j - tp_i \geq dur_{min}$  then accept  $P^*(i, j)$ .
    if  $dur_{min} \leq tp_j - tp_i + trav(d, bl_i) \leq dur_{max}$  then accept  $s + P^*(i, j)$ .
    if  $dur_{min} \leq tp_j - tp_i + trav(el_j, d) \leq dur_{max}$  then accept  $P^*(i, j) + t$ .
    if  $dur_{min} \leq tp_j - tp_i + trav(d, bl_i) + trav(el_j, d) \leq dur_{max}$ 
        then accept  $s + P^*(i, j) + t$ .
end.
    
```

 Figure 4. Procedure RETRIEVEPIECES(i, k)

The term $s + P^*(i, j)$ denotes the extension of path $P^*(i, j)$ by adding s as the initial node, and the term $P^*(i, j) + t$ denotes the extension of path $P^*(i, j)$ by adding t as the terminal node. Accepted pieces are stored in a set of potential pieces to be used for the duty generation phase. The running time of procedure RETRIEVEPIECES(i, k) is $O(n)$, where n equals the number of trips. Thus, the overall running time of the piece generation algorithm is $O(nm)$.

When generating sets of pieces we have to assure that the column generation optimality condition is satisfied, that is, no negative reduced cost duties exist at convergence of the column generation procedure. To facilitate the discussion, we do not consider s and t as relief points. Hence, the first relief point of a path s, i, \dots, j is the point corresponding to node i . Let S denote the set of pieces containing all pairs of shortest paths generated by the algorithm in Figure 3. Furthermore, we assume that the duty generation algorithm satisfies the column generation optimality condition, that is, negative reduced cost duties are detected as long as corresponding paths exist in the duty generation network. Recall that we have assumed that the feasibility of a piece of work depends on the starting and ending relief point only. We now have the following proposition.

Proposition 1. If the duty generation algorithm with set S as the set of pieces does not return a negative reduced cost duty under the assumptions stated above, then there are no duties with negative, reduced costs at all.

Proof. Consider an arbitrary feasible duty with pieces not in S , and suppose that this duty has negative reduced cost. Any piece in this duty that is not in S can be replaced by a piece in S with the same starting and ending relief point and with lower or equal reduced cost, while the duty remains feasible due to the assumptions stated above. Thus, if at least one duty with negative reduced costs exist, then at least one such duty exists with pieces in S only. Hence, if no negative reduced cost duties are found with pieces in S only, then no negative reduced cost duties exist when considering all pieces. ■

7. COMPLETE INTEGRATION: NO CHANGEOVERS

7.1. Model and algorithm

The formulation presented in the preceding section is valid irrespective of the fact whether or not changeovers are allowed. It only differs in the set of feasible duties K . However, if changeovers are not allowed, we can also use a different formulation. Recall that a changeover is the change of vehicle of a driver during his break. We define a *combined duty* as a feasible vehicle duty *and* one or more corresponding feasible crew duties. The compatibility of vehicle and crew duties is guaranteed by the definition of a combined duty. Therefore, the VCSP without changeovers can be modelled in a straightforward way as a set covering problem, to which we will refer as VCSP2.

The algorithm for VCSP2 is almost similar to the one given in Figure 2 for VCSP1. One difference is that we have to take into account that no changeovers are allowed in the initial solution. We do this by solving the CSP where the set of duties consists only of duties without changeovers. Since model VCSP2 is a pure set covering model, Lagrangian relaxations can be obtained in a similar way as discussed in Section 5.2. For the model VCSP2 we must generate combined duties instead of duties. How this can be done, will be discussed in detail in Section 7.2. The other difference with Figure 2 is that we obtain a feasible solution in a similar way as for CSP by solving a set covering problem with the columns generated during the lower bound phase. So the difference is that we do not compute a feasible vehicle schedule first, since combined duties already define a vehicle schedule in itself.

7.2. The column generation subproblem

The column generation pricing problem for the integrated model without changeovers needs an additional procedure in order to generate combined duties using previously generated crew duties as input (see Section 6.3). This results in a three-phase procedure. While the piece and duty generation networks are created only once, the combined duty generation network needs to be built every time combined duties are generated, using the duties generated previously.

Figure 5 demonstrates the procedures.

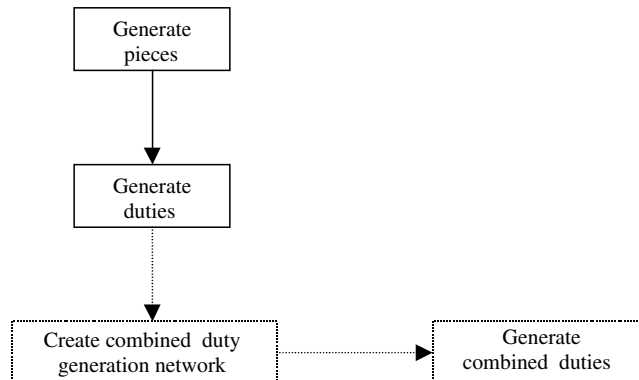


Figure 5. Column generation problem

A disadvantage of the three-phase procedure is that it is possible to generate a duty twice in the second phase, and to generate a combined duty twice in the third phase. This problem is already mentioned in Section 6.3.

We will discuss the differences in the generation of pieces compared to the integrated model with changeovers and describe the generation of combined duties.

7.2.1. Generation of pieces of work

We can use the same network for generation of pieces as in Section 6.4, but the costs in this network are different. Let λ_p be the Lagrangian multiplier corresponding to a trip or trip task, and associated with the set covering constraints in model VCSP2. Then the reduced cost is defined as

$$\bar{b}_{ij} = \begin{cases} b_{ij} - \lambda_p, & \text{for each arc } (i,j) \text{ corresponding to a trip or trip task } p. \\ b_{ij}, & \text{otherwise.} \end{cases}$$

Of course, we can still use the same all-pairs shortest path algorithm to generate the set of pieces.

7.2.2. Generation of combined duties

An additional procedure is necessary to generate combined duties for model VCSP2. The input for this procedure is a set of duties generated by the two phase procedure, that is, pieces of work are generated as described in Section 6.4, while duties are generated as described in Section 5.3. The purpose of the third phase is to find combined duties with negative reduced cost. This is done by constructing a combined duty generation network based on the given duties, and solving a sequence of (unconstrained) shortest path problems in this network. The negative reduced cost combined duties generated in this way are added to the master problem, while convergence of the column generation algorithm is attained once no negative reduced cost combined duties are obtained.

Combined duty generation network G is constructed as follows. Nodes correspond to the set of duties previously generated plus a source and a sink. An arc from the source to a duty node exists if the starting relief point of the duty is located at the depot, while an arc from a duty to the sink exists if the ending relief point of the duty is located at the depot. Furthermore, an arc between two duty nodes i and j exists if duty i and j form a compatible pair of duties on one vehicle, that is, if $el_i = bl_j$ and $et_i \leq bt_j$. In addition, an arc exists if two duties can serve together on one vehicle while interrupting each other. Suppose for example that two duties i and j each contain two pieces p_i, q_i and p_j, q_j , respectively. Then, arc (i,j) exists if pairs (p_i, p_j) , (p_j, q_i) , and (q_i, q_j) form compatible pairs of pieces.

Reduced vehicle trip costs and vehicle deadhead costs are incorporated in the reduced cost of pieces and duties. This is achieved by adjusting arc costs in the piece generation network. The costs of arc (i,j) , $i \neq s$ in network G equals the reduced cost of duty j plus the vehicle cost of the corresponding deadhead. If $i = s$, then the fixed vehicle cost is added as well. A path starting at s and ending at t is denoted by an (s,t) -path. We can determine the combined duty corresponding to the shortest (s,t) -path in network G in $O(m)$ time, where m is the number of arcs. This shortest path is added to the master problem if its cost is negative, while the column

generation procedure is terminated otherwise. In order to speed up column generation by trying to generate more than one column per iteration, we solve several shortest path problems by leaving out nodes and arcs of paths already found. For more details, we refer to Freling (1997).

If the set S of pieces is generated as described previously, that is, all-pairs shortest paths are included, and all feasible duties that can be constructed using the pieces in S are input for the combined duty generation, then optimality can be guaranteed, at convergence of the column generation procedure. This can be proven in a similar way as Proposition 1.

8. INDEPENDENT CREW SCHEDULING

It is an obvious observation that the specification of vehicle schedules will put certain constraints on the crew schedules (and vice versa). Because vehicles are often much more flexible to schedule than crews, it may be inefficient to schedule vehicles without considering crew scheduling. We can measure the potential benefit of integration with respect to cost efficiency by comparing the cost of the traditional sequential approach, that is, first solve the SDVSP and then the CSP, with a lower bound on the optimal value of the VCSP. To this end, we define the ICSP as follows: Given a set of trip tasks corresponding to a set of trips, and given the traveling times between each pair of locations, find a minimum cost crew schedule such that all trip tasks are covered in exactly one duty and all duties satisfy crew feasibility constraints. Note that vehicles are completely ignored in this problem. Clearly, a lower bound can be computed by independently solving the SDVSP and the ICSP and adding the respective optimal values.

If the cost of the sequential approach is much higher than the lower bound, it may be that the crew scheduling solution will improve significantly when considering integration as compared to the sequential approach. On the other hand, we know that there is no need to integrate if the cost of the sequential solution is close to the lower bound. Note that, because the vehicle schedules are identical for the sequential and independent approach, the potential benefit can be measured by comparing the cost of the crew schedules only.

For the ICSP we propose the same formulation as for the CSP, but the difference is that the set of possible duties is much larger, because the vehicle schedule is not given in advance. Furthermore, we know that some duties will never appear in the optimal solution. For example, consider two duties i and j which are equivalent except that j contains an extra dh-task from the depot to the first relief point. Then, duty j will never appear in an optimal solution if $d_i < d_j$. In general, if the cost of a duty increases with the duration, in an optimal solution no duty begins or ends at the depot. So vehicle movements to and from the depot are not covered. In particular, only duties beginning and ending with a trip task need to be considered.

The algorithm to solve the ICSP is similar to the one of VCSP2 (see Section 7.1), but instead of combined duties we have only duties.

9. COMPUTATIONAL TESTS FOR BUS AND DRIVER SCHEDULING

In this section we present a computational study of integrated bus and driver scheduling. The aim of this study is to investigate the effectiveness of integration as compared to the traditional sequential approach, and the applicability of the proposed techniques in practice.

9.1. Data and parameter settings

We have used data from the RET, the public transport company in Rotterdam, the Netherlands, that provides passenger service by bus, metro, and tram. The data corresponds to bus and driver scheduling of individual bus lines. Some constraints at RET are very complex and particular for this company only. Therefore, we have changed some RET specific constraints into more general constraints. In another paper Freling, Huisman, and Wagelmans (2001), we considered the real RET restrictions, which will make the problem much more difficult. An example of such a constraint is the round-trip-condition, which limits the total layover (nondriving) time in a round-trip (two consecutive trips such that a driver drives from location A via location B back to location A).

The restrictions that we have taken into account, are as follows. A driver can only be relieved by another driver at the start or end of a trip or at the depot and continuous attendance is required. This implies that a trip always corresponds to exactly one trip task and the number of relief points is equal to twice the number of trips. If a driver starts/ends his duty at the depot, there is a sign-on/sign-off time of 11 and 12 min, respectively. Another restriction is that every bus waits at least 3 min after a trip at the end location before starting the next trip. There are four different types of duties and their properties are given in Figure 6.

The following points are relevant for all the different approaches proposed in this section:

1. The primary objective is to minimize the sum of buses and drivers used in the schedule. For bus scheduling we have a secondary objective of minimizing a combination of travel and idle time as defined in Section 6.1. For crew scheduling we do not have a secondary objective. We use fixed costs in order to assure that the minimum number of buses and drivers is the primary objective. For the computational results presented in this section, we only report the number of buses and drivers since these are relevant for our study.
2. The pricing problems are solved separately for each type of duty. We work with reduced networks for each type of duty, that is, only nodes corresponding to pieces of work are considered which could be part of a duty of a certain type.
3. The Lagrangian multipliers are initialized to zero. The maximum number of iterations k_{\max} is initially set to 100 for problem instances with a small number of trips; for instances with a larger number of trips, we take for the independent approach $k_{\max} = 1000$. In every iteration of the column generation algorithm, after solving the master problem, k_{\max} is increased by 3. A greedy heuristic is used to get only nonnegative Lagrangian cost variables in the master problem. The initial values of the Lagrangian multipliers for a new

	Type 1 (tripper)		Type 2 (early duty)		Type 3 (normal duty)		Type 4 (late duty)	
	Minimum	Maximum	Minimum	Maximum	Minimum	Maximum	Minimum	Maximum
Number of pieces	1	1	2	2	2	2	2	2
Start time (min)				645	646	914	915	
End time (min)				1,177		1,334		
Piece length (min)	30	300	30	300	30	240	30	300
Break length (min)			15	90	15	90	15	90
Duty length (min)	30	300		532		420		570
Work time (min)	30	300		532		420		570

Figure 6. Properties of all different types of duties

master problem are set to the values resulting from the previous master problem, before the greedy heuristic was applied.

4. The column generation is stopped once the difference between the estimate of the lower bound for the overall problem and the current lower bound is less than 1% (0.1% for VCSP1) or if no improvement in the lower bound is obtained for five iterations (250 iterations for VCSP1). The latter is the so-called *tailing-off criterion*.
5. For the integrated approach with changeovers allowed, we generate 10 feasible solutions, where 5 follow from the last iterations of the column generation and the other 5 follow from the last iterations of a subgradient optimization after the convergence of the column generation step. We choose the best feasible solution. If the gap between the lower bound and the best feasible solution is large, we can easily adapt our approach to compute some additional feasible solutions in the last subgradient optimization, that is, the one after the convergence of the column generation step.
6. The strategies for generating columns in the pricing problem have been obtained after extensive testing and tuning on different types of problems (see Freling (1997)). We add for all types of duties the first 200 duties with negative reduced cost that we find. This means that we do not take the duties with the most negative reduced cost. The advantage of this strategy is that we do not have to generate all the duties in all iterations. The disadvantage is that the convergence process is slower, but overall this is faster than the alternative strategy.

All tests are executed on a Pentium 400 PC with 128 MB of computer memory.

We cannot compare our test results with the results of others, because—to the best of our knowledge—only Haase, Desaulniers, and Desrosiers (2001) perform a computational study for simultaneous scheduling of vehicles and crews, but the characteristics of their random data and our real life data differ in the following four respects:

1. Their trip lengths are much longer than ours; on average they have between 2 and 3 trips per duty and about 4 per vehicle, whereas we have, for the largest problem, on average more than 10 trips per duty and 25 per vehicle.
2. They consider a network of four lines, while we consider individual bus lines; note that this means that for the same total number of trips, the (average) frequency per line is higher for our problem instances.
3. They assume that bus drivers are only allowed to start and end their duty at the depot, whereas in our case drivers can start and end at every relief point (i.e., at the start/end of every trip and at the depot).
4. They have only one duty type with two pieces and one consisting of a single piece, while we consider three duty types with two pieces and one with a single piece.

9.2. Results of different approaches

In this section we show the results of the sequential, independent, and integrated approach. For the sequential and the integrated approach we look at the cases with and without changeovers.

9.2.1. Sequential approach

In Figure 7 we present the computational results for the sequential method with and without changeovers. The number of trips for prob24 is equal to 24 and so on. Since the primary

	Changeovers						No changeovers					
	Prob24	Prob42	Prob72	Prob113	Prob148	Prob238	Prob24	Prob42	Prob72	Prob113	Prob148	Prob238
Lower	12	19	15	23	34	27	12	19	15	26	45	30
Upper	12	19	15	24	34	29	12	20	15	26	45	33
Buses	6	9	5	8	11	9	6	9	5	8	11	9
Drivers	6	10	10	16	23	20	6	11	10	18	34	24
Gap (%)	0.00	0.00	0.00	4.17	0.00	6.90	0.00	5.00	0.00	0.00	0.00	9.09

Figure 7. Results sequential approach RET data

	Prob24	Prob42	Prob72	Prob113	Prob148	Prob238
Nodes	48	84	144	226	296	476
Arcs	132	312	460	846	1,376	2,260
Iterations	5	6	9	9	15	20
Average pieces	463	1,568	4,606	9,891	16,237	52,538
Duties	529	1,689	6,026	9,329	28,208	70,307
CPU m.	0	0	3	7	1,261	1,845
CPU p.	1	8	41	136	6,907	14,422
CPU t.	1	8	44	178	8,169	16,274
Lower	6	9	9	13	19	16
Upper	6	9	9	14	22	18
Gap (%)	0.00	0.00	0.00	7.14	13.64	11.11
CPU	0	0	0	50	53	52

Figure 8. Results independent approach RET data

objective is to minimize the sum of buses and drivers, we give the lower and upper bounds in terms of this sum.

Because the initial step of the integrated approach is the sequential approach, the difference in the computation time between the sequential and integrated approach is equal to the computation time of the lower and upper bound phase of the latter approach (which we will present later on). Moreover, the computation times of the sequential approach is negligible. For these reasons, we have not mentioned them here.

9.2.2. Independent approach

In Figure 8 we show the computational results for the independent crew scheduling. The top part of the figure focuses on the lower bound phase and the bottom part on the upper bound phase. For the lower bound phase, we report the number of nodes and arcs in the piece generation network G' , the number of column generation iterations, the average number of pieces of work during the column generation, the number of duties in the final master problem, the total computation times in seconds required for the master problems, the pricing problems, and the lower bound phase, and the resulting lower bound on the number of drivers in the solution, respectively. For the upper bound phase, we report the cost (number of drivers) of the feasible solution, the relative gap between lower and upper bounds, and the computation times in seconds for the upper bound phase.

As can be seen from Figure 8, the gap between the lower and upper bounds for the independent crew scheduling is zero for smaller problem instances, but it increases for the larger instances prob113 (7.14%), prob148 (13.64%), and prob238 (11.11%). As mentioned before, for

the larger problem instances we perform more iterations of the subgradient optimization than for the smaller instances. By doing so, we obtain improved lower bounds for these problems. For example, for prob148 we improved the lower bound from 17 to 19. This is possible, because only a small part of the computational effort is in the master problem and the major part is in the pricing problem. In general, such large gaps in a column generation approach occur due to not generating columns during the search for integer solutions.

Recall from Section 8 that we can measure the potential benefit of integration by comparing solutions obtained by the sequential approach with solutions obtained by the independent approach. The minimum number of drivers resulting from model ICSP is a lower bound on the minimum number of drivers overall. Thus, if we compare the number of drivers resulting from the independent approach with the number of drivers in the CSP solution, taking the gap with the lower bound into account, we also have an indication of the potential benefit. That is, we have an indication of how many drivers could be saved from the sequential approach by an integrated one.

9.2.3. Integrated approach

In Figure 9 we show the computational results for the integrated approach with and without changeovers. We report the same information as in Figure 8, except that we also report the number of combined duties in the case of “no changeovers” and the number of buses in the solution.

The gap between the lower and upper bounds is at most 3.57% for the approach with changeovers and 7.32% if changeovers are not allowed. However, we should note that we do not have a guaranteed lower bound in all cases, because of the tailing-off criterion. For prob72 we have computed an additional feasible solution, as described in the Section 9.1. If we would not have done this, the best solution was 15 and the gap 6.67%. In the case where changeovers are not allowed we have computed the feasible solution for prob148 and prob238 with integer programming solver of CPLEX instead of the heuristic of Caprara, Fischetti, and Toth (1999), because for these problems the heuristic performed poorly. By doing this, we saved one vehicle and three drivers for prob148 and one vehicle and one driver for prob238.

	Changeovers						No changeovers					
	Prob24	Prob42	Prob72	Prob113	Prob148	Prob238	Prob24	Prob42	Prob72	Prob113	Prob148	Prob238
Nodes	48	84	144	226	296	476	48	84	144	226	296	476
Arcs	132	312	460	846	1,376	2,260	132	312	460	846	1,376	2,260
Iterations	5	12	19	44	106	254	3	4	12	9	15	14
Average pieces	463	1,568	4,606	9,891	16,237	52,538	1,367	4,706	8,716	19,511	33,435	101,187
Duties	512	1,711	7,063	12,644	29,667	80,945	2,201	3,780	5,092	3,150	8,984	16,536
Combined duties							213	438	628	323	1,816	4,972
CPU m.	0	2	18	150	911	27,644	0	0	7	0	3	12
CPU p.	1	14	34	837	4,188	214,436	1	19	289	41	255	3,927
CPU t.	1	18	56	1,024	5,185	242,080	1	19	296	41	258	3,942
Real_lower	12	18	14	23	33	27	12	18	14	26	38	30
Upper	12	18	14	23	34	28	12	18	14	26	41	32
Buses	6	9	5	8	11	9	6	9	5	8	12	9
Drivers	6	9	9	15	23	19	6	9	9	18	29	23
Gap (%)	0.00	0.00	0.00	0.00	2.94	3.57	0.00	0.00	0.00	0.00	7.32	6.25
CPU	4	3	37	212	351	1,907	0	0	42	0	29	35

Figure 9. Results integrated approach on RET data

Problem	Independent				Changeovers								No changeovers							
					Sequential				Integration				Sequential				Integration			
	Low	Upp	Bus	Driv	Low	Upp	Bus	Driv	Low	Upp	Bus	Driv	Low	Upp	Bus	Driv	Low	Upp	Bus	Driv
Prob24	12	12	6		12	12	6	6	12	12	6	6	12	12	6	6	12	12	6	6
Prob42	18	18	9		19	19	9	10	18	18	9	9	19	20	9	11	18	18	9	9
Prob72	14	14	9		15	15	5	10	14	14	5	9	15	15	5	10	14	14	5	9
Prob113	21	22	14		23	24	8	16	23	23	8	15	26	26	8	18	26	26	8	18
Prob148	30	33	22		34	34	11	23	33	34	11	23	45	45	11	34	38	41	12	29
Prob238	25	27	18		27	29	9	20	27	28	9	19	30	33	9	24	30	32	9	23

Figure 10. Comparison between different approaches for RET data

Again, the major computational effort is in the pricing problem. For prob148 (with, changeovers), the computation time of the lower bound phase is about 86 min. About 80% is spent on the pricing problem, where an average number of 16,237 pieces of work and a total number of 29,667 duties are generated. In the next section we give a comparison between the different approaches.

9.3. A comparison of different approaches

In this section, we are interested in a comparison of the approaches discussed in this paper. In Figure 10 we present the results obtained by the sequential, the independent and the integrated approach for the RET data. For the sequential approach, the upper bounds (lower bounds) are obtained by the summation of the minimum number of vehicles obtained by solving the SDVSP and the (lower bound on the) number of drivers resulting from the Lagrangian heuristic for the CSP.

The number of buses is not mentioned for the independent approach because it is the same as for the sequential approach. The results for the integrated approach correspond to the best lower and upper bounds obtained using different column generation strategies (as mentioned in the previous subsection). For the case with changeovers, comparing the sequential with the integrated approach shows that the savings when integrating bus and driver scheduling are at most one driver. This confirms the expectation based on the potential savings resulting from comparing the sequential with the independent approach. In four out of six test cases, the integrated solution was better than the sequential solution due to the saving of one driver. However, for the case without changeovers a significant saving is obtained with the integrated approach. This also confirms the expectation based on the potential savings resulting from comparing the sequential with the independent approach. In four out of six test cases, the integrated solution was better than the sequential solution due to the saving of up to five drivers at the cost of only one additional bus. Interestingly, for the small problems the integrated approach gives the same results for the cases with and without changeovers, but for the larger problems allowing changeovers can save many drivers and buses.

10. CONCLUSIONS

The results reported in the previous section show that we can get good solutions within reasonable computation times on a personal computer. We did not expect that it would be useful to integrate under every circumstance. Based on the results we can conclude that the

benefit of integration may be significant when changeovers are not allowed. In practice, it often occurs that either changeovers are not possible due to long distances or changeovers are not allowed for legal or technical reasons. When changeovers are allowed in our test problems, one driver may be saved by considering an integrated approach. Of course, the saving of one driver is considerable when taking into account that we investigated instances for one bus line each. The interpretation of the computational results depends on the ratio between the fixed vehicle and crew costs. If fixed vehicle costs are much higher as compared to crew costs it becomes less attractive to apply the integrated approach. On the other hand, if crew costs are higher the integrated approach becomes more attractive.

Furthermore, we have seen that the results of the independent approach give a good indication of the potential benefit of integration.

Finally, we would like to remark the following. Although, from an employment point of view, it may be interpreted negatively when we say that drivers can be saved by using integration, a different point of view is that the cost savings may allow for an increase in service while no driver needs to lose his or her job.

ACKNOWLEDGMENTS

The authors are grateful to the RET, the Rotterdam public transport company, for providing the data. They would also like to thank an associate editor and an anonymous referee for their comments.

REFERENCES

- Ball, M., L. Bodin, and R. Dial, "A matching based heuristic for scheduling mass transit crews and vehicles," *Transport. Sci.*, **17**, 4–31 (1983).
- Barnhart, C., E. L. Johnson, G. L. Nemhauser, M. W. P. Savelsbergh, and P. H. Vance, "Branch-and-price: Column generation for solving huge integer programs," *Oper. Res.*, **46**, 316–329 (1998).
- Beasley, J. E., "Lagrangian relaxation," in C. R. Reeves (ed.), *Modern Heuristic Techniques for Combinatorial Problems*, McGraw-Hill, London, 1995, pp. 243–303.
- Bodin, L., B. Golden, A. Assad, and M. Ball, "Routing and scheduling of vehicles and crews: The state of the art," *Comput. Oper. Res.*, **10**(2), 63–211 (1983).
- Caprara, A., M. Fischetti, and P. Toth, "A heuristic algorithm for the set covering problem," *Oper. Res.*, **47**, 730–743 (1999).
- Carraraesi, P., L. Girardi, and M. Nonato, "Network models, Lagrangean relaxation and subgradients bundle approach in crew scheduling problems," in J. R. Daduna, I. Branco, J. M. Pinto Paixão (eds.), *Computer-Aided Transit Scheduling: Proceedings of the Sixth International Workshop*, Springer Verlag, Berlin, 1995, pp. 188–212.
- Darby-Dowman, K., J. K. Jachnik, R. L. Lewis, and G. Mitra, "Integrated decision support systems for urban transport scheduling: Discussion of implementation and experience," in J. R. Daduna and A. Wren (eds.), *Computer-Aided Transit Scheduling: Proceedings of the Fourth International Workshop*, Springer Verlag, Berlin, 1988, pp. 226–239.
- Desrochers, M., "Shortest path problems with resource constraints," *Technical Report GERAD G-88-27*, École des Hautes Études Commerciales, Montréal, 1988.
- Desrochers, M. and F. Soumis, "A column generation approach to the urban transit crew scheduling problem," *Transport. Sci.*, **23**, 1–13 (1989).
- Falkner, J. C. and D. M. Ryan, "Express: Set partitioning for bus crew scheduling in Christchurch," in M. Desrochers, J. M. Rousseau (eds.), *Computer-Aided Scheduling: Proceedings of the Fifth International Workshop*, Springer Verlag, Berlin, 1992, pp. 359–378.
- Fisher, M. L., "The Lagrangian relaxation method for solving integer programming problems," *Manag. Sci.*, **27**, 1–18 (1981).
- Freling, R., "Models and Techniques for Integrating Vehicle and Crew Scheduling," Tinbergen Institute, Erasmus University Rotterdam, 1997.
- Freling, R., C. G. E. Boender, and J. M. Pinto Paixão, "An integrated approach to vehicle and crew scheduling," *Technical Report 9503/A*, Econometric Institute, Erasmus University Rotterdam, Rotterdam, 1995.

- Freling, R., D. Huisman, and A. P. M. Wagelmans, "Applying an integrated approach to vehicle and crew scheduling in practice," in S. Voß and J. R. Daduna (eds.), *Computer-Aided Scheduling of Public Transport*, Springer, Berlin, 2001, pp. 73–90.
- Freling, R., J. M. Pinto Paixão, and A. P. M. Wagelmans, "Models and algorithms for vehicle scheduling," *Transport. Sci.*, **35**, 165–180 (2001).
- Freling, R., A. P. M. Wagelmans, and J. M. Pinto Paixão, "An overview of models and techniques for integrating vehicle and crew scheduling," in N. H. M. Wilson (ed.), *Computer-Aided Transit Scheduling*, Springer Verlag, Berlin, 1999, pp. 441–460.
- Gaffi, A. and M. Nonato, "An integrated approach to extra-urban crew and vehicle scheduling," in N. H. M. Wilson (ed.), *Computer-Aided Transit Scheduling*, Springer Verlag, Berlin, 1999, pp. 103–128.
- Geoffrion, A. M., "Lagrangian relaxation for integer programming," *Math. Program. Study*, **2**, 82–114 (1974).
- Haase, K. and C. Friberg, "An exact branch and cut algorithm for the vehicle and crew scheduling problem," in N. H. M. Wilson (ed.), *Computer-Aided Transit Scheduling*, Springer Verlag, Berlin, 1999, pp. 63–80.
- Haase, K., G. Desaulniers and J. Desrosiers, "Simultaneous vehicle and crew scheduling in urban mass transit systems," *Transport. Sci.*, **35**(3), 286–303 (2001).
- Huisman, D., "A note on model VCSP1," World Wide Web, <http://www.few.eur.nl/few/people/huisman/publications.htm>, 2000.
- Klabjan, D., E. L. Johnson, G. L. Nemhauser, E. Gelman, and S. Ramaswamy, "Airline crew scheduling with time windows and plane count constraints," *Transport. Sci.*, **36**, 337–348 (2002).
- Patrikalakis, I. and D. Xerocostas, "A new decomposition scheme of the urban public transport scheduling problem," in M. Desrochers and J. M. Rousseau (eds.), *Computer-Aided Transit Scheduling: Proceedings of the Fifth International Workshop*, Springer Verlag, Berlin, 1992, pp. 407–425.
- Rousseau, J.M. and J. Y. Blais, "Hastus: An interactive system for buses and crew scheduling," in J. M. Rousseau (ed.), *Computer Scheduling of Public Transport 2*, North Holland, Amsterdam, 1985, pp. 45–60.
- Scott, D., "A large linear programming approach to the public transport scheduling and cost model," in J. M. Rousseau (ed.), *Computer Scheduling of Public Transport 2*, North Holland, Amsterdam 1985, pp. 473–491.
- Tosini, E. and C. Vercellis, "An interactive system for extra-urban vehicle and crew scheduling problems," in J. R. Daduna and A. Wren (eds.), *Computer-Aided Transit Scheduling: Proceedings of the Fourth International Workshop*, Springer Verlag, Berlin, 1988, pp. 41–53.